

SESSION 8

VECTORISATION & DATA LAYOUT

MASSIMILIANO FASI



Durham
University

Scalar and vector operations

$$z \leftarrow x + y$$

Scalar operation

$$\begin{bmatrix} z_0 \\ \vdots \\ z_n \end{bmatrix} \leftarrow \begin{bmatrix} x_0 \\ \vdots \\ x_n \end{bmatrix} + \begin{bmatrix} y_0 \\ \vdots \\ y_n \end{bmatrix}$$

Vector operation

Two realizations

- ▶ lockstepping (GPUs SIMT)
- ▶ large vector registers (x86 extensions)

Large vector registers



+



=



Large vector registers



+




=




+



=




Large vector registers



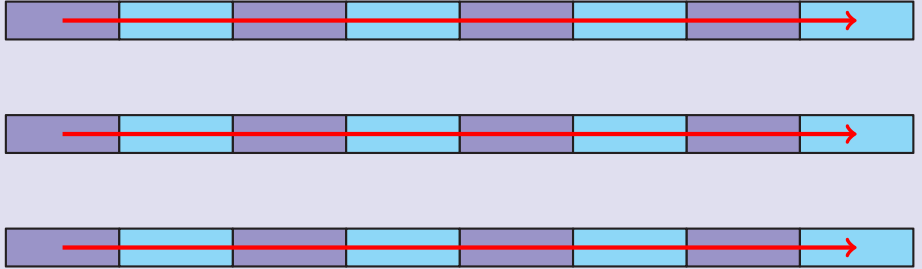
► SIMD lanes

Large vector registers



- ▶ SIMD lanes
- ▶ vertical operation

Large vector registers



- ▶ SIMD lanes
- ▶ vertical operation
- ▶ horizontal operation

Vector extensions

Arch.	Extension	bits	binary32	binary64
x86	SSE	128	4	–
x86	SSE2	128	4	2
x86	AVX	256	8	4
x86	AVX2 (FMA)	256	8	4
x86	AVX512	512	16	8
ARM	SVE	128–2048	4–64	2–32

- ▶ **SSE** Streaming SIMD Extension
- ▶ **AVX** Advanced Vector eXtension
- ▶ **SVE** Scalable Vector Extension

Compiler x86 options



Architectures

- ▶ `-march=x86-64`
- ▶ `-march=core-avx2`
- ▶ `-march=skylake-avx512`
- ▶ `-march=znver2`
- ▶ `-march=native`

Extensions

- ▶ `-mmmx`
- ▶ `-msse`
- ▶ `-msse4.2`
- ▶ `-mavx2`
- ▶ `-mavx512f` (Foundation)

The GCC flag `--help=target` shows **all** target-specific options.

VECTORISING C AND C++ CODE

Vectorisation in practice

1. Automatic optimisation

▶ `g++ -fopt-info`



▶ `icpc -qopt-report`



2. Compiler loop-specific `#pragma` directives¹

3. OpenMP² vectorisation `#pragma` directives

4. Compiler built-in (intrinsic) functions



5. Hand-written assembly

¹ Pragmas are used to give additional information to the compiler.

² Open Multi-Processing.

Unrolling a for loop

```
for (int i = 0; i < N; ++i)  
    a[i] = b[i] + c[i];
```

Unrolling a for loop

```
for (int i = 0; i < N; ++i)
    a[i] = b[i] + c[i];
```

```
for (int i = 0; i < 4 * (N / 4); i += 4) {
    a[i]      = b[i]      + c[i];
    a[i+1]    = b[i+1]    + c[i+1];
    a[i+2]    = b[i+2]    + c[i+2];
    a[i+3]    = b[i+3]    + c[i+3];
}
for (; i < N; ++i)
    a[i]      = b[i]      + c[i];
```

Requirements for automatic vectorisation

1. iteration count known beforehand
2. no jumps (`break/continue`)
3. no exceptions
4. no loop carried dependency
5. no nested loops
6. no if statements (almost)
7. no function calls (almost)

This requires `-ftree-vectorize` (included in `-O3`).



Will the compiler vectorise this?

```
double A* = (double *) malloc(N * N * sizeof *A);  
double B* = (double *) malloc(N * N * sizeof *B);  
double C* = (double *) malloc(N * N * sizeof *C);
```

```
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        C[i*N + j] += A[i*N + j] * B[i*N + j];
```

Will the compiler vectorise this?

```
double A* = (double *) malloc(N * N * sizeof *A);  
double B* = (double *) malloc(N * N * sizeof *B);  
double C* = (double *) malloc(N * N * sizeof *C);
```

```
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        C[i*N + j] += A[i*N + j] * B[i*N + j];
```

```
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        C[j*N + i] += A[j*N + i] * B[j*N + i];
```


Will the compiler vectorise this?

```
double A* = (double *) malloc(N * sizeof *A);
```

```
for (int i = 0; i != N; ) {  
    tmp = N;  
    N = A[i];  
    A[i] = tmp;  
}
```

Will the compiler vectorise this?

```
double A* = (double *) malloc(N * sizeof *A);
```

```
for (int i = 0; i != N; ) {  
    tmp = N;  
    N = A[i];  
    A[i] = tmp;  
}
```

```
for (int i = 0; i <= N; i++) {  
    if (A[i] > 0)  
        sum += A[i];  
}
```

Will the compiler vectorise this?

```
double A* = (double *) malloc(N * sizeof *A);
```

```
for (int i = 2; i < N; i++)  
    A[i] = (A[i-1] + A[i-2]) / 2;
```

Will the compiler vectorise this?

```
double A* = (double *) malloc(N * sizeof *A);
```

```
for (int i = 2; i < N; i++)  
    A[i] = (A[i-1] + A[i-2]) / 2;
```

```
void foo(double *A, double *B, double *C, int N)  
{  
    for (int i = 0; i < N; i++)  
        A[i] = (B[i] + C[i]) / 2;  
}
```

What are B and C?


Pointer aliasing (C/C++)

```
void foo(double *A, double *B, double *C, int N)
{
    for (int i = 0; i < N; i++)
        A[i] = (B[i] + C[i]) / 2;
}
```

```
void bar(double *c, int n)
{
    double *a = c + 2;
    double *b = c + 1;
    foo(a, b, c, n-3);
}
```

Pointer aliasing (C/C++)

```
void bar(double *c, int n)
{
    double *a = c + 2;
    double *b = c + 1;
    foo(a, b, c, n-3);
}
```



Solutions to pointer aliasing

1. Some compilers can handle it on their own
 - ▶ Multiple versions of the loop are generated
 - ▶ Run-time check for aliasing
 - ▶ Appropriate version is used

Solutions to pointer aliasing

1. Some compilers can handle it on their own
 - ▶ Multiple versions of the loop are generated
 - ▶ Run-time check for aliasing
 - ▶ Appropriate version is used
2. Tell the compiler with `-fno-alias`

Solutions to pointer aliasing

1. Some compilers can handle it on their own
 - ▶ Multiple versions of the loop are generated
 - ▶ Run-time check for aliasing
 - ▶ Appropriate version is used
2. Tell the compiler with `-fno-alias`
3. We can guarantee pointers will not alias
 - ▶ `double * restrict` (in C99 or newer)
 - ▶ `double * __restrict__` (in C++)

Compiler loop-specific `#pragma` directives

```
#pragma <directive>\n<for_loop>
```

Ignore vector dependencies

► `g++` : `#pragma GCC ivdep`



► `icpc` : `#pragma ivdep`



Force loop unrolling factor

► `g++` : `#pragma GCC unroll(<factor>)`



► `icpc` : `#pragma unroll(<factor>)`



OpenMP vectorisation `#pragma` directives



```
#pragma omp simd [<clause>[[,<clause>]]...]\n<for_loop>
```

For **vertical** operations, <clause> can be

- ▶ `safelen(<length>)`: unrolling factor safe to use.
- ▶ `simdlen(<length>)`: number of SIMD lanes to use.
- ▶ `linear(<list>[:<step>])`: step for variables in <list>.
- ▶ `if([simd :] <expr>)`: vectorise only if <expr> is true.
- ▶ `collapse(<num>)`: collapse <num> levels of nested loops.

Reduction

For **horizontal** operations, <clause> can be

```
reduction([ < modifier > ,] < identifier > : < list > )
```

where <identifier> can be

- ▶ an arithmetic operation: +, *, -, max, min
- ▶ a logical or bitwise operation: &, &&, |, ||, ^

and <list> is a list of variables. For <modifier>s, see .

Vectorised functions

```
#pragma omp declare simd [<clause>[[,<clause>]...]\n[#pragma omp declare simd [<clause>[[,<clause>]\n[...]\n<function_definition_or_declaration>
```

- ▶ Generates multiple (vectorised) versions of the function.
- ▶ **However**, compilers will often inline, then vectorise.
- ▶ Use `-fno-inline` (g++) or `-qno-inline` (icpc) to check.

Some examples

- ▶ Ignoring vector dependencies
- ▶ Safe forward dependencies
- ▶ Reduction in inner product
- ▶ Declare SIMD function



Some examples

- ▶ Ignoring vector dependencies
- ▶ Safe forward dependencies
- ▶ Reduction in inner product
- ▶ Declare SIMD function



USEFUL REMINDERS

- ▶ Don't forget `-fopenmp` (g++) or `-qopenmp` (icpc)!
- ▶ `-fopt-info` (g++) and `-qopt-report` (icpc) can help.
- ▶ Memory **must be contiguous**

