

Session 5: Cache blocking/tiling

COMP52315: performance engineering

Lawrence Mitchell*

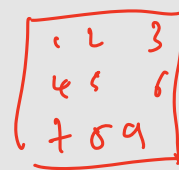
*`lawrence.mitchell@durham.ac.uk`

An exemplar problem

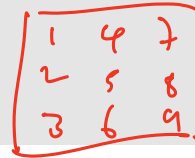
Matrix transpose

$$B_{ij} \leftarrow A_{ji} \quad A, B \in \mathbb{R}^{n \times n}$$

```
double *a, *b;  
...  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        b[i*N + j] = a[j*N + i];
```



1	2	3
4	5	6
7	8	9



1	4	7
2	5	8
3	6	9

So far, we've talked about how to measure performance, and perhaps determine that it is bad.

⇒ what can we do about it?

→ expect memory
bus-limited.

- 0 flushing past ops.
- 1 load of an entry of A
- 1 store of B
- every loop iteration.

Matrix transpose: simple performance model

Set up our expectation

- N^2 loads, N^2 stores, no compute
- ⇒ all we're doing is copying data
- Hence we might expect to see performance close to that of the streaming memory bandwidth, independent of matrix size.

Matrix transpose: simple performance model

Set up our expectation

- N^2 loads, N^2 stores, no compute
- ⇒ all we're doing is copying data
- Hence we might expect to see performance close to that of the streaming memory bandwidth, independent of matrix size.

	Matrix size	BW [GByte/s]
L_2 in cache ←	128×128	22
L_3 cache ← {	256×256	13
	512×512	13
main memory ← {	1024×1024	5
	2048×2048	1.6
	4096×4096	0.9

On laptop
with
stream BW
 $A \sim 10$ Gb/s.



What went wrong?

```
double *a, *b;  
...  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        b[i*N + j] = a[j*N + i];
```


- We have streaming access to **b**, but stride- N access to **a**.
- If both matrices fit in cache, this is OK, and a reasonable model of time is $T_{\text{cache}} = N^2(t_{\text{read}} + t_{\text{write}})$.
- Note that the reads of **a** load a full cache line, but use only 8 bytes of it.
- Better model $T_{\text{mem}} = N^2(8t_{\text{read}} + t_{\text{write}})$

A picture

load that provides
a new cache line to
be brought in cache
these are a
cache hit.



These are
not accessed
in context.



if there are too
many rows.
⇒ switch
before we
run.

Cache locality

- Since we have strided access to \mathbf{a} , we need to hold LN bytes in the cache to get any reuse, where L is the cache line size in. This is not possible for large matrices.
- A mechanism to fix this is to *reorder* the loop iterations to preserve spatial locality.

of cache lines.

Idea

- Break loop iteration space into blocks
 - *strip-mining*
 - *loop reordering*

Strip mining

- Break a loop into blocks of consecutive elements

Before

```
for ( int i = 0; i < N; i++ )  
    a[i] = f(i);
```

After

```
for ( int ii = 0; ii < N; ii += stride)  
    for ( int i = ii; i < min(N, ii + stride); i++)  
        a[i] = f(i);
```

- Not that useful for just a single loop, although there are circumstances where one might use it \longrightarrow compilers do this for vectorization.

Strip mining multiple loops

- Let's do the same for both loops of the transpose:

Before

```
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
        a[i*N + j] = a[j*N + i];
```

After

```
for (int ii = 0; ii < N; ii += stridei)  
    for (int i = ii; i < min(N, ii+stridei); i++)  
        for (int jj = 0; jj < N; jj += stridej)  
            for (int j = jj; j < min(N, jj+stridej); j++)  
                b[i*N + j] = a[j*N + i];
```

- Haven't yet made any change to the performance


Reorder loops

After permuting i and jj loops

```
for (int ii = 0; ii < N; ii += stridei)
    for (int jj = 0; jj < N; jj += stridej)
        for (int i = ii; i < min(N, ii+stridei); i++)
            for (int j = jj; j < min(N, jj+stridej); j++)
                b[i*N + j] = a[j*N + i];
```

- Two free parameters `stridei` and `stridej`
- Need to choose these appropriately to levels in the cache hierarchy
- Ideally block for L1, L2, L3, etc...
- The extra logic adds some overhead

small enough to fit in cache.



Why is it “tiling”?

Iteration over B .

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63

Why is it “tiling”?

Iteration over A.

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63

0	1	2	3	4	5	6	7
8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23
24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55
56	57	58	59	60	61	62	63

Does it work?

- Have a go, I provide some sample code for which you can tune the blocking parameters.

⇒ Exercise 7.

....	1100
...-	1101
-..	1110
-...-	1111

4 bit cache 1's
index:

⇒ 16 cache lines.

A second problem

Matrix-Matrix multiplication

$$C_{ij} \leftarrow C_{ij} + \sum_k A_{ik} B_{kj} \quad A, B, C \in \mathbb{R}^{n \times n}$$

```
for (int i = 0; i < n; i++)  
  for (int j = 0; j < n; j++)  
    for (int k = 0; k < n; k++)  
      C[i*n + j] += A[i*n + k] * B[k*n + j];
```

Perfect for modern hardware

Perfect with balance

Same story here (or at least it was in the 90s!).

always in a pair
⇒ FMA.

$\leftarrow \begin{matrix} N^3 & \text{mults} \\ N^3 & \text{adds} \end{matrix}$

⇒ Touch $O(N^2)$ data.
⇒ should be flop limited.

(Another) simple model for computation

- Simple model of memory, two levels: “fast” and “slow”
- Initially all data in slow memory

m number of data elements moved between fast and slow memory

t_m time per slow memory operation

f number of flops

$t_f \ll t_m$ time per flop

$q =: f/m$ average flops per slow memory access

inspired by a 1981 paper.

- Minimum time to solution (all data in fast memory)

- Typical time


compute
 \downarrow
 $ft_f + mt_m = ft_f \left(1 + \frac{t_m}{t_f} \frac{1}{q} \right)$
memory movement hardware
algorithm

- t_m/t_f property of hardware, q property of algorithm

Naïve matrix-multiply

```
for (int i = 0; i < n; i++)  
  for (int j = 0; j < n; j++)  
    for (int k = 0; k < n; k++)  
      C[i*n + j] = C[i*n + j] + A[i*n + k] * B[k*n + j];
```


- Algorithm does $2n^3 = \mathcal{O}(n^3)$ flops and touches $3 \cdot 8n^2$ bytes of memory
- q potentially $\mathcal{O}(n)$, arbitrarily large for large n .



Naïve matrix-multiply

```
for (int i = 0; i < n; i++)  
    // Read row i of A into fast memory  
    for (int j = 0; j < n; j++)  
        // Read Cij into fast memory  
        // Read column j of B into fast memory  
        for (int k = 0; k < n; k++)  
            C[i*n + j] = C[i*n + j] + A[i*n + k] * B[k*n + j];  
        // Write Cij back to slow memory
```

everything we need
is in fast memory.



Naïve matrix-multiply

Number of slow memory references

$m = n^3$ each column of B is read n times

$+ n^2$ each row of A is read ~~once~~ once *$\leftarrow n$ rows of $A \Rightarrow n^2$ data*


$+ 2n^2$ each entry of C is read once and written once

$$= (n^3 + 3n^2)$$

Hence

$$\lim_{n \rightarrow \infty} q = \frac{f}{m} = \frac{2n^3}{(n^3 + 3n^2)} = 2$$

*So this algorithm is bad.
 $t_f + \left(1 + \frac{t_m}{t_s} \frac{1}{2}\right)$*



From model to prediction

- So for a triply-nested loop structure, the *best* time to solution our model predicts is:

$$T = t_f f \left(1 + \frac{t_m}{2t_f} \right)$$

- Recall that on modern hardware, memory *latency* is around 200 cycles per cache line. So let's approximate $t_m \approx 200/8 = 25$, and say $t_f = 1$.


$$T = t_f f (1 + 25/2) = 13.5 t_f f$$

- Maximally 7% peak.
- This is *only* an estimate.

large N limit
takes ~ 13.5 times
as long as
peak flop rate.

Measurement

- Single core Intel i5-8259U.
 - 2 4-wide FMAs per cycle \Rightarrow 16 DP FLOPs/cycle.
- \Rightarrow Peak is $3.6 \cdot 16 = 57.6$ GFLOPs/s, model predicts 4.03GFLOPs/s.



How to improve reuse?

- Problem is that we move rows and columns into fast memory, and then evict them
- Need way of keeping the loaded data in fast memory as long as possible.

⇒ tile iterations

```
// Treat  $A, B, C \in (\mathbb{R}^{b \times b})^{N \times N}$   
// that is,  $N \times N$  matrices where each entry is a  $b \times b$  matrix.  
for (int i = 0; i < N; i++)  
  for (int j = 0; j < N; j++)  
    // Read block  $C_{ij}$  into fast memory ←  
    for (int k = 0; k < n; k++)  
      // Read block  $A_{ik}$  into fast memory  
      // Read block  $B_{kj}$  into fast memory  
      // Do matrix multiply on the blocks  
       $C[i*N + j] = C[i*N + j] + A[i*N + k] * B[k*N + j];$  ←  
    // Write block  $C_{ij}$  back to slow memory
```

~ block multiply ~
→ hiding three more loops.

How to improve reuse?

- Problem is that we move rows and columns into fast memory, and then evict them
- Need way of keeping the loaded data in fast memory as long as possible.

⇒ tile iterations

```
// Treat  $A, B, C \in (\mathbb{R}^{b \times b})^{N \times N}$   
// that is,  $N \times N$  matrices where each entry is a  $b \times b$  matrix.  
for (int ii = 0; ii < N; ii++)  
    for (int jj = 0; jj < N; jj++)  
        for (int kk = 0; kk < N; kk++)  
            for (int i_ = 0; i_ < b; i_++)  
                for (int j_ = 0; j_ < b; j_++)  
                    for (int k_ = 0; k_ < b; k_++) {  
                        const int i = ii*b + i_;  
                        const int j = jj*b + j_;  
                        const int k = kk*b + k_;  
                        C[i*n + j] = C[i*n + j] + A[i*n + k] * B[k*n + j];  
                    }  
}
```


What did that do to the data movement?

$$N < n.$$

before this was n^3

$$\begin{aligned} m &= Nn^2 \quad \text{each block of } B \text{ is read } N^3 \text{ times} \Rightarrow N^3 b^2 = N^3 (n/N)^2 = Nn^2 \\ &+ Nn^2 \quad \text{each block of } A \text{ is read } N^3 \text{ times} \rightarrow \text{before this was } n^2 \\ &+ 2n^2 \quad \text{each block of } C \text{ is read once and written once} \\ &= 2n^2(N + 1) \end{aligned}$$

\rightarrow Before it was $n^3 + n^2$.

Hence

$$\lim_{n \rightarrow \infty} q = \frac{f}{m} = \frac{2n^3}{2n^2(N + 1)} = \frac{n}{N} = b$$

- $b \gg 2$ so much better than previously. Can improve performance by increasing b as long as blocks still fit in fast memory!
- Detailed analysis of blocked algorithms in Lam, Rothberg, and Wolf *The Cache Performance and Optimization of Blocked Algorithms* (1991)

From model to machine characteristics

- Arbitrarily choose a “fast” algorithm to be $\geq 50\%$ peak, this requires

$$ft_f \left(1 + \frac{t_m}{t_f} \frac{1}{q} \right) \leq 2t_f f \Leftrightarrow \frac{t_m}{t_f} \frac{1}{q} \leq 1 \Leftrightarrow q \geq \frac{t_m}{t_f}$$

- Again, approximate $t_m = 25$, $t_f = 1$

$\Rightarrow b \approx q \geq 25$.

- Need to hold all three $b \times b$ matrices in cache

\Rightarrow Need space for $3b^2 = 3 \cdot 25^2 = 1875$ matrix *entries*, approximately 14.6KB of fast memory M_{fast} .

- This is smaller than L1, but larger than fits in registers.

\Rightarrow should be achievable.

Is this the best we can do?

Theorem

Hong and Kung (1981) Any reorganization of this algorithm that only exploits associativity has

$$q = \mathcal{O}(\sqrt{M_{fast}})$$

and the number of data elements moved between slow and fast memory is

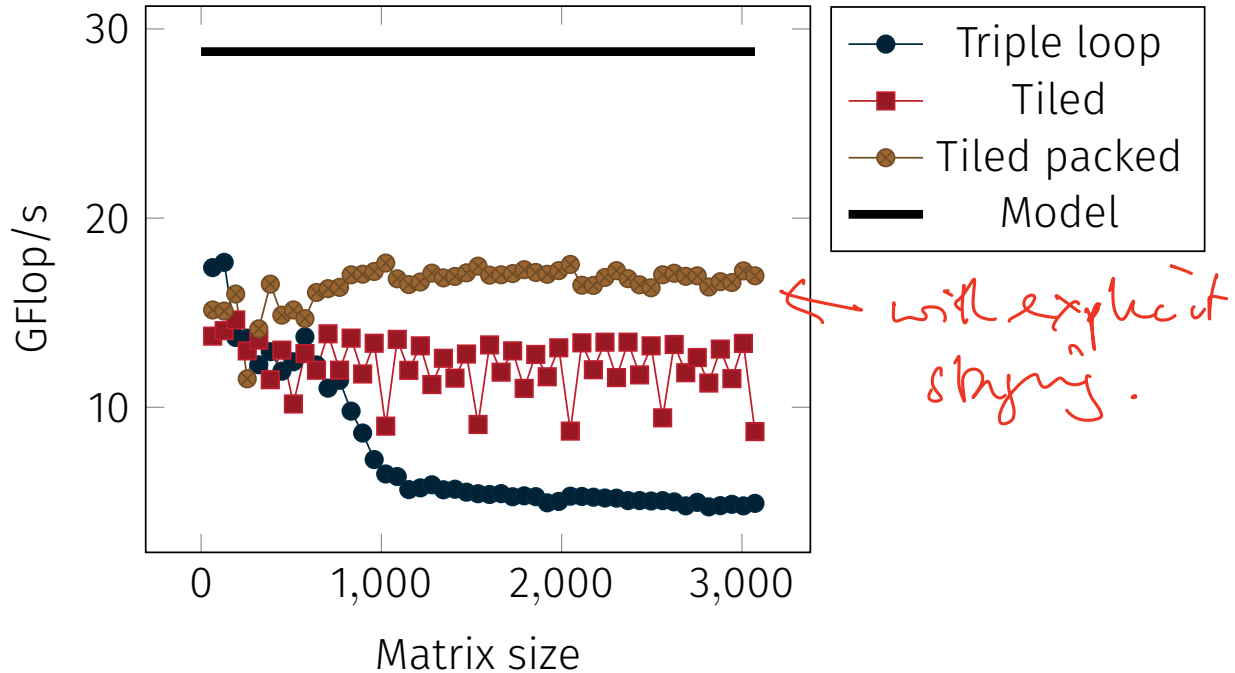
$$\Omega\left(\frac{n^3}{\sqrt{M_{fast}}}\right)$$

- Exact values for the bounds are not known, the best bounds are provided by Smith and van de Geijn (2017) **arXiv: 1702.02017 [cs.CC]**
- The GotoBLAS/OpenBLAS approach approaches these bounds.


Matching reality with models

- I provide some sample code that implements this scheme
⇒ Exercise 8.

Is this the best we can do?



Is this the best we can do?



What accounts for this difference?

- Managed to get big matrices to behave like small ones with naive code.
⇒ reaching in-cache performance of the starting point.
- For better results, need to
 1. Block for registers and all levels of cache → 4 levels of memory system.
 2. Perform data-layout transformation to promote (better) vectorisation
- Will look more at data layout transforms next time.

Summary

- Loop tiling can *significantly* improve performance of nested loops.
- Particularly important to exploit data reuse.
- For the “last mile” we have to do more. Mostly the same idea, but thinking hard about data layout and explicit vectorisation.
- Simple models can be used to motivate whether things are worth trying.